

dlnd_face_generation

July 31, 2018

1 Face Generation

In this project, you'll use generative adversarial networks to generate new images of faces. ###
Get the Data You'll be using two datasets in this project: - MNIST - CelebA

Since the celebA dataset is complex and you're doing GANs in a project for the first time, we want you to test your neural network on MNIST before CelebA. Running the GANs on MNIST will allow you to see how well your model trains sooner.

If you're using [FloydHub](#), set `data_dir` to `"/input"` and use the [FloydHub data ID](#) `"R5KrjnANiKVhLWApXhNBe"`.

```
In [1]: data_dir = '/data'
        !pip install matplotlib==2.0.2
        # FloydHub - Use with data ID "R5KrjnANiKVhLWApXhNBe"
        #data_dir = '/input'

        """
        DON'T MODIFY ANYTHING IN THIS CELL
        """

        import helper

        helper.download_extract('mnist', data_dir)
        helper.download_extract('celeba', data_dir)
```

Collecting matplotlib==2.0.2

Downloading https://files.pythonhosted.org/packages/60/d4/6b6d8a7a6bc69a1602ab372f6fc6e88ef8/

100% || 14.6MB 51kB/s eta 0:00:01

Requirement already satisfied: pyparsing!=2.0.0,!=2.0.4,!=2.1.2,!=2.1.6,>=1.5.6 in /opt/conda/

Requirement already satisfied: six>=1.10 in /opt/conda/lib/python3.6/site-packages (from matpl

Requirement already satisfied: cyclor>=0.10 in /opt/conda/lib/python3.6/site-packages/cyclor-0

Requirement already satisfied: numpy>=1.7.1 in /opt/conda/lib/python3.6/site-packages (from ma

Requirement already satisfied: python-dateutil in /opt/conda/lib/python3.6/site-packages (from

Requirement already satisfied: pytz in /opt/conda/lib/python3.6/site-packages (from matplotlib

Installing collected packages: matplotlib

Found existing installation: matplotlib 2.1.0

Uninstalling matplotlib-2.1.0:

Successfully uninstalled matplotlib-2.1.0

Successfully installed matplotlib-2.0.2

You are using pip version 9.0.1, however version 18.0 is available.You should consider upgrading via the `python -m pip install --upgrade pip` command.

Found mnist Data

Found celeba Data

1.1 Explore the Data

1.1.1 MNIST

As you're aware, the [MNIST](#) dataset contains images of handwritten digits. You can view the first number of examples by changing `show_n_images`.


```
In [2]: show_n_images = 25
```

```
"""
DON'T MODIFY ANYTHING IN THIS CELL
"""

%matplotlib inline
%config InlineBackend.figure_format = 'retina'
import os
from glob import glob
from matplotlib import pyplot

mnist_images = helper.get_batch(glob(os.path.join(data_dir, 'mnist/*.jpg'))[:show_n_images])
pyplot.imshow(helper.images_square_grid(mnist_images, 'L'), cmap='gray')
```

```
Out[2]: <matplotlib.image.AxesImage at 0x7fda75efd048>
```




1.1.2 CelebA

The [CelebFaces Attributes Dataset \(CelebA\)](#) dataset contains over 200,000 celebrity images with annotations. Since you're going to be generating faces, you won't need the annotations. You can view the first number of examples by changing `show_n_images`.

```
In [3]: show_n_images = 25
```

```
"""  
DON'T MODIFY ANYTHING IN THIS CELL  
"""  
mnist_images = helper.get_batch(glob(os.path.join(data_dir, 'img_align_celeba/*.jpg'))  
                                pyplot.imshow(helper.images_square_grid(mnist_images, 'RGB'))
```

```
Out [3]: <matplotlib.image.AxesImage at 0x7fda74cf6ef0>
```



1.2 Preprocess the Data

Since the project's main focus is on building the GANs, we'll preprocess the data for you. The values of the MNIST and CelebA dataset will be in the range of -0.5 to 0.5 of 28x28 dimensional images. The CelebA images will be cropped to remove parts of the image that don't include a face, then resized down to 28x28.

The MNIST images are black and white images with a single [color channel]([https://en.wikipedia.org/wiki/Channel_\(digital_image%29\)](https://en.wikipedia.org/wiki/Channel_(digital_image%29))) while the CelebA images have [3 color channels (RGB color channel)]([https://en.wikipedia.org/wiki/Channel_\(digital_image%29#RGB_Images\)](https://en.wikipedia.org/wiki/Channel_(digital_image%29#RGB_Images))). ## Build

the Neural Network You'll build the components necessary to build a GANs by implementing the following functions below: - model_inputs - discriminator - generator - model_loss - model_opt - train

1.2.1 Check the Version of TensorFlow and Access to GPU

This will check to make sure you have the correct version of TensorFlow and access to a GPU

```
In [4]: """
        DON'T MODIFY ANYTHING IN THIS CELL
        """

        from distutils.version import LooseVersion
        import warnings
        import tensorflow as tf

        # Check TensorFlow Version
        assert LooseVersion(tf.__version__) >= LooseVersion('1.0'), 'Please use TensorFlow version 1.0 or higher'
        print('TensorFlow Version: {}'.format(tf.__version__))

        # Check for a GPU
        if not tf.test.gpu_device_name():
            warnings.warn('No GPU found. Please use a GPU to train your neural network.')
        else:
            print('Default GPU Device: {}'.format(tf.test.gpu_device_name()))
```

```
TensorFlow Version: 1.3.0
Default GPU Device: /gpu:0
```

1.2.2 Input

Implement the model_inputs function to create TF Placeholders for the Neural Network. It should create the following placeholders: - Real input images placeholder with rank 4 using image_width, image_height, and image_channels. - Z input placeholder with rank 2 using z_dim. - Learning rate placeholder with rank 0.

Return the placeholders in the following the tuple (tensor of real input images, tensor of z data)

```
In [5]: import problem_unittests as tests

        def model_inputs(image_width, image_height, image_channels, z_dim):
            """
            Create the model inputs
            :param image_width: The input image width
            :param image_height: The input image height
            :param image_channels: The number of image channels
            :param z_dim: The dimension of Z
            :return: Tuple of (tensor of real input images, tensor of z data, learning rate)
            """
```

```

# TODO: Implement Function
input_real = tf.placeholder(tf.float32, (None, image_height, image_width, image_channels))
input_z = tf.placeholder(tf.float32, (None, z_dim), name='input_z')
learning_rate = tf.placeholder(tf.float32, (), name='learning_rate')

return input_real, input_z, learning_rate

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""
tests.test_model_inputs(model_inputs)

```

Tests Passed

1.2.3 Discriminator

Implement discriminator to create a discriminator neural network that discriminates on images. This function should be able to reuse the variables in the neural network. Use `tf.variable_scope` with a scope name of "discriminator" to allow the variables to be reused. The function should return a tuple of (tensor output of the discriminator, tensor logits of the discriminator).

```

In [75]: def discriminator(images, reuse=False):
    """
    Create the discriminator network
    :param images: Tensor of input image(s)
    :param reuse: Boolean if the weights should be reused
    :return: Tuple of (tensor output of the discriminator, tensor logits of the discriminator)
    """

    # Alpha value for computing Leaky ReLU activations
    alpha = 0.2

    # Drop rate for dropout layers (layers 1 and 3 out of
    # four total layers)
    drop_rate = 0.3

    # TODO: Implement Function
    with tf.variable_scope('discriminator', reuse=reuse):

        # Inputs will be of size 28 x 28 x dim_output_channel.
        # Do not use batch normalization on first conv layer, in order
        # to avoid introducing weird artifacts into the images.
        layer_1 = tf.layers.conv2d(images, filters=64, kernel_size=(5,5), strides=(2,2))
        # Leaky ReLU activation
        lrelu1 = tf.maximum(alpha * layer_1, layer_1)
        lrelu1 = tf.layers.dropout(lrelu1, rate=drop_rate)

```

```

# Shape now 14x14x64

layer_2 = tf.layers.conv2d(lrelu1, filters=128, kernel_size=(5,5), strides=(2
# Batch normalize all subsequent layers, save
# for the output layer. We are only interested in using the
# discriminator to help train the generator, so training
# parameter is set to True when calling tf.layers.batch_normalization.
batch_norm2 = tf.layers.batch_normalization(layer_2, training=True)
# Leaky ReLU activation
lrelu2 = tf.maximum((alpha * batch_norm2), batch_norm2)
# Shape now 7x7x128

layer_3 = tf.layers.conv2d(lrelu2, filters=256, kernel_size=(5,5), strides=(2
# Batch normalize
batch_norm3 = tf.layers.batch_normalization(layer_3, training=True)
# Leaky ReLU activation
lrelu3 = tf.maximum((alpha * batch_norm3), batch_norm3)
lrelu3 = tf.layers.dropout(lrelu3, rate=drop_rate)
# Shape now 4x4x256

layer_4 = tf.layers.conv2d(lrelu3, filters=512, kernel_size=(5,5), strides=(2
# Batch normalize
batch_norm4 = tf.layers.batch_normalization(layer_4, training=True)
# Leaky ReLU activation
lrelu4 = tf.maximum((alpha * batch_norm4), batch_norm4)
# Shape now 2x2x512

# Flatten final convolutional layer into a fully
# connected layer. Number of units in flat layer
# is equal to number of units in final convolutional
# layer (lrelu5).
fully_connected = tf.reshape(lrelu4, [-1, 2*2*512])

# Only one unit in logits because we are using
# a sigmoid function.
logits = tf.layers.dense(inputs=fully_connected, units=1, activation=None)
output = tf.sigmoid(logits)

return output, logits

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""
tests.test_discriminator(discriminator, tf)

```

Tests Passed

1.2.4 Generator

Implement generator to generate an image using z . This function should be able to reuse the variables in the neural network. Use `tf.variable_scope` with a scope name of "generator" to allow the variables to be reused. The function should return the generated $28 \times 28 \times \text{out_channel_dim}$ images.

```
In [76]: def generator(z, out_channel_dim, is_train=True):
        """
        Create the generator network
        :param z: Input z
        :param out_channel_dim: The number of channels in the output image
        :param is_train: Boolean if generator is being used for training
        :return: The tensor output of the generator
        """

        # Alpha value for computing Leaky ReLU activations
        alpha = 0.2

        # Re-use generator variables when not training.
        if is_train:
            reuse = False
        else:
            reuse = True

        # TODO: Implement Function
        with tf.variable_scope('generator', reuse=reuse):

            # Fully connected input layer
            fc_input_layer = tf.layers.dense(z, 4*4*512, activation=None)

            # Reshape to start the convolutional stack
            layer1 = tf.reshape(fc_input_layer, [-1, 4, 4, 512])
            # batch normalize
            batch_norm1 = tf.layers.batch_normalization(inputs=layer1, training=is_train)
            # Leaky ReLU activation
            lrelu1 = tf.maximum((alpha * batch_norm1), batch_norm1)
            # Shape now 4x4x512

            # Initialize with Xavier initialization, use kernel size
            # of 4, and padding of 'valid' in order to resize shape
            # from 4x4x512 to 7x7x256.
            layer2 = tf.layers.conv2d_transpose(lrelu1, filters=256, kernel_size=(4,4), s
            # batch normalize
            batch_norm2 = tf.layers.batch_normalization(inputs=layer2, training=is_train)
            # Leaky ReLU activation
            lrelu2 = tf.maximum((alpha * batch_norm2), batch_norm2)
            # Shape now 7x7x256
```

```

layer3 = tf.layers.conv2d_transpose(lrelu2, filters=128, kernel_size=(5,5), s
# batch normalize
batch_norm3 = tf.layers.batch_normalization(inputs=layer3, training=is_train)
# Leaky ReLU activation
lrelu3 = tf.maximum((alpha * batch_norm3), batch_norm3)
# Shape now 14x14x128

layer4 = tf.layers.conv2d_transpose(lrelu3, filters=64, kernel_size=(5,5), st
# batch normalize
batch_norm4 = tf.layers.batch_normalization(inputs=layer4, training=is_train)
# Leaky ReLU activation
lrelu4 = tf.maximum((alpha * batch_norm4), batch_norm4)
# Shape now 28x28x64

# The output layer. No batch normalization nor activation
# function applied here.
logits = tf.layers.conv2d_transpose(lrelu4, filters=out_channel_dim, kernel_s
# Shape now is 28 x 28 x out_channel_dim

output = tf.tanh(logits)

return output

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""
tests.test_generator(generator, tf)

```

Tests Passed

1.2.5 Loss

Implement `model_loss` to build the GANs for training and calculate the loss. The function should return a tuple of (discriminator loss, generator loss). Use the following functions you implemented: `- discriminator(images, reuse=False) - generator(z, out_channel_dim, is_train=True)`

```

In [77]: def model_loss(input_real, input_z, out_channel_dim):
        """
        Get the loss for the discriminator and generator
        :param input_real: Images from the real dataset
        :param input_z: Z input
        :param out_channel_dim: The number of channels in the output image
        :return: A tuple of (discriminator loss, generator loss)
        """
        # TODO: Implement Function

```



```

# Generator model output
g_model_image_output = generator(input_z, out_channel_dim, is_train=True)

# Discriminator model outputs and logits for real and
# generated images.
d_model_real_image_output, d_logits_real = discriminator(input_real, reuse=False)
d_model_fake_image_output, d_logits_fake = discriminator(g_model_image_output, reuse=False)

# Discriminator loss
d_loss_real = tf.reduce_mean(
    tf.nn.sigmoid_cross_entropy_with_logits(logits=d_logits_real, labels=tf.ones_like(d_logits_real))
)
d_loss_fake = tf.reduce_mean(
    tf.nn.sigmoid_cross_entropy_with_logits(logits=d_logits_fake, labels=tf.zeros_like(d_logits_fake))
)
d_loss = d_loss_real + d_loss_fake

# Generator loss
g_loss = tf.reduce_mean(
    tf.nn.sigmoid_cross_entropy_with_logits(logits=d_logits_fake, labels=tf.ones_like(d_logits_fake))
)

return d_loss, g_loss

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""
tests.test_model_loss(model_loss)

```

Tests Passed

1.2.6 Optimization

Implement `model_opt` to create the optimization operations for the GANs. Use `tf.trainable_variables` to get all the trainable variables. Filter the variables with names that are in the discriminator and generator scope names. The function should return a tuple of (discriminator training operation, generator training operation).

```

In [81]: def model_opt(d_loss, g_loss, learning_rate, beta1):
        """
        Get optimization operations
        :param d_loss: Discriminator loss Tensor
        :param g_loss: Generator loss Tensor
        :param learning_rate: Learning Rate Placeholder
        :param beta1: The exponential decay rate for the 1st moment in the optimizer
        :return: A tuple of (discriminator training operation, generator training operation)
        """
        # TODO: Implement Function

```

```

# Get all trainable variables for both the discriminator
# and generator.
t_vars = tf.trainable_variables()
d_vars = [var for var in t_vars if var.name.startswith('discriminator')]
g_vars = [var for var in t_vars if var.name.startswith('generator')]

# Optimize
# Start with tf.control_dependencies since we're using
# batch normalization.
with tf.control_dependencies(tf.get_collection(tf.GraphKeys.UPDATE_OPS)):
    d_train_opt = tf.train.AdamOptimizer(learning_rate, beta1=beta1).minimize(d_loss)
    g_train_opt = tf.train.AdamOptimizer(learning_rate, beta1=beta1).minimize(g_loss)

return d_train_opt, g_train_opt

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""
tests.test_model_opt(model_opt, tf)

```

Tests Passed

1.3 Neural Network Training

1.3.1 Show Output

Use this function to show the current output of the generator during training. It will help you determine how well the GANs is training.

```

In [82]: """
DON'T MODIFY ANYTHING IN THIS CELL
"""

import numpy as np

def show_generator_output(sess, n_images, input_z, out_channel_dim, image_mode):
    """
    Show example output for the generator
    :param sess: TensorFlow session
    :param n_images: Number of Images to display
    :param input_z: Input Z Tensor
    :param out_channel_dim: The number of channels in the output image
    :param image_mode: The mode to use for images ("RGB" or "L")
    """
    cmap = None if image_mode == 'RGB' else 'gray'
    z_dim = input_z.get_shape().as_list()[-1]

```

```

example_z = np.random.uniform(-1, 1, size=[n_images, z_dim])

samples = sess.run(
    generator(input_z, out_channel_dim, False),
    feed_dict={input_z: example_z})

images_grid = helper.images_square_grid(samples, image_mode)
pyplot.imshow(images_grid, cmap=cmap)
pyplot.show()

```

1.3.2 Train

Implement train to build and train the GANs. Use the following functions you implemented:

- model_inputs(image_width, image_height, image_channels, z_dim)
- model_loss(input_real, input_z, out_channel_dim)
- model_opt(d_loss, g_loss, learning_rate, beta1)

Use the show_generator_output to show generator output while you train. Running show_generator_output for every batch will drastically increase training time and increase the size of the notebook. It's recommended to print the generator output every 100 batches.

In [86]: *# How often to print out discriminator and generator loss*

```
print_every = 100
```

```
# How often we display image samples created by the generator.
```

```
show_every = 500
```

```
# To plot the discriminator and generator losses at the end
# of training.
```

```
losses = []
```

```
def train(epoch_count, batch_size, z_dim, learning_rate, beta1, get_batches, data_shape)
```

```
    """
```

```
    Train the GAN
```

```
    :param epoch_count: Number of epochs
```

```
    :param batch_size: Batch Size
```

```
    :param z_dim: Z dimension
```

```
    :param learning_rate: Learning Rate
```

```
    :param beta1: The exponential decay rate for the 1st moment in the optimizer
```

```
    :param get_batches: Function to get batches
```

```
    :param data_shape: Shape of the data
```

```
    :param data_image_mode: The image mode to use for images ("RGB" or "L")
```

```
    """
```

```
    # TODO: Build Model
```

```
    image_height = data_shape[1]
```

```
    image_width = data_shape[2]
```

```
    image_channels = data_shape[3]
```

```
    # Create input and learning rate placeholders
```

```

input_real, input_z, lr = model_inputs(image_width, image_height, image_channels,

# Set the learning rate
lr = learning_rate

# Get the losses
d_loss, g_loss = model_loss(input_real, input_z, image_channels)

# Get the optimizers
d_train_opt, g_train_opt = model_opt(d_loss, g_loss, lr, beta1)

# To count the number of batches trained on so far.
steps = 0

with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    for epoch_i in range(epoch_count):
        for batch_images in get_batches(batch_size):

            # TODO: Train Model

            # Increase the iteration count by one each batch.
            steps += 1

            # Sample random noise for Generator input
            batch_z = np.random.uniform(-1, 1, size=(batch_size, z_dim))

            # Real input images are in range [-0.5, 0.5]. They need to
            # be scaled up to the range [-1,1]. This is because the
            # generator will output images scaled between [-1,1].
            batch_images = batch_images * 2

            # Run the optimizers
            _ = sess.run(d_train_opt, feed_dict={input_real: batch_images, input_z: batch_z})
            _ = sess.run(g_train_opt, feed_dict={input_real: batch_images, input_z: batch_z})

            # Print out loss and sample images created by the
            # generator every print_every batches.
            if steps % print_every == 0:
                # Get the losses and print them out
                train_loss_d = d_loss.eval({input_z: batch_z, input_real: batch_images})
                train_loss_g = g_loss.eval({input_z: batch_z})

                # Save losses to view after training
                losses.append((train_loss_d, train_loss_g))

                print("Epoch {} of {}      ".format(epoch_i+1, epochs),
                      "Steps: {}      ".format(steps),

```

```

        "Discriminator Loss: {:.4f}    ".format(train_loss_d),
        "Generator Loss: {:.4f}".format(train_loss_g))

    # Print out sample images created by the
    # generator every show_every batches.
    if steps % show_every == 0:
        print("Sample Output:")

        # Display a sample of images that the generator is
        # capable of producing at this point.
        show_generator_output(sess, n_images=25, input_z=input_z, out_channel_dim=out_channel_dim)

    # Print out losses and sample outputs at the
    # end of each epoch, as well:
    # Get the losses and print them out
    train_loss_d = d_loss.eval({input_z: batch_z, input_real: batch_images})
    train_loss_g = g_loss.eval({input_z: batch_z})

    print("End of Epoch {} ".format(epoch_i+1),
          "Steps: {}    ".format(steps),
          "Discriminator Loss: {:.4f}    ".format(train_loss_d),
          "Generator Loss: {:.4f}\n".format(train_loss_g),
          "Sample Output:")

    # Display a sample of images that the generator is
    # capable of producing at this point.
    show_generator_output(sess, n_images=25, input_z=input_z, out_channel_dim=out_channel_dim)

```

1.3.3 MNIST

Test your GANs architecture on MNIST. After 2 epochs, the GANs should be able to generate images that look like handwritten digits. Make sure the loss of the generator is lower than the loss of the discriminator or close to 0.

```

In [84]: batch_size = 32
        z_dim = 100
        learning_rate = 0.002
        beta1 = 0.5

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """

        epochs = 2


        mnist_dataset = helper.Dataset('mnist', glob(os.path.join(data_dir, 'mnist/*.jpg')))
        with tf.Graph().as_default():

```

```
train(epochs, batch_size, z_dim, learning_rate, beta1, mnist_dataset.get_batches,
      mnist_dataset.shape, mnist_dataset.image_mode)
```


Epoch 1 of 2	Steps: 100	Discriminator Loss: 0.0130	Generator Loss: 4.6008
Epoch 1 of 2	Steps: 200	Discriminator Loss: 0.8191	Generator Loss: 0.6961
Epoch 1 of 2	Steps: 300	Discriminator Loss: 1.3351	Generator Loss: 0.3463
Epoch 1 of 2	Steps: 400	Discriminator Loss: 1.5166	Generator Loss: 0.4212
Epoch 1 of 2	Steps: 500	Discriminator Loss: 1.1836	Generator Loss: 0.5099

Sample Output:




Epoch 1 of 2	Steps: 600	Discriminator Loss: 0.3584	Generator Loss: 1.7919
Epoch 1 of 2	Steps: 700	Discriminator Loss: 0.5225	Generator Loss: 1.3143
Epoch 1 of 2	Steps: 800	Discriminator Loss: 2.3893	Generator Loss: 0.1455
Epoch 1 of 2	Steps: 900	Discriminator Loss: 1.0032	Generator Loss: 1.8950
Epoch 1 of 2	Steps: 1000	Discriminator Loss: 1.4977	Generator Loss: 0.3494

Sample Output:




Epoch 1 of 2	Steps: 1100	Discriminator Loss: 0.9067	Generator Loss: 0.9648
Epoch 1 of 2	Steps: 1200	Discriminator Loss: 1.2889	Generator Loss: 0.5191
Epoch 1 of 2	Steps: 1300	Discriminator Loss: 2.0054	Generator Loss: 0.2097
Epoch 1 of 2	Steps: 1400	Discriminator Loss: 1.8005	Generator Loss: 1.3637
Epoch 1 of 2	Steps: 1500	Discriminator Loss: 0.4730	Generator Loss: 1.6247

Sample Output:




Epoch 1 of 2	Steps: 1600	Discriminator Loss: 1.7917	Generator Loss: 0.2148
Epoch 1 of 2	Steps: 1700	Discriminator Loss: 0.4747	Generator Loss: 1.5208
Epoch 1 of 2	Steps: 1800	Discriminator Loss: 2.2938	Generator Loss: 0.2072
End of Epoch 1	Steps: 1875	Discriminator Loss: 1.0904	Generator Loss: 0.7474

Sample Output:




Epoch 2 of 2	Steps: 1900	Discriminator Loss: 1.1383	Generator Loss: 0.6202
Epoch 2 of 2	Steps: 2000	Discriminator Loss: 1.1646	Generator Loss: 0.6200

Sample Output:




Epoch 2 of 2	Steps: 2100	Discriminator Loss: 0.5850	Generator Loss: 1.6146
Epoch 2 of 2	Steps: 2200	Discriminator Loss: 1.3990	Generator Loss: 0.4078
Epoch 2 of 2	Steps: 2300	Discriminator Loss: 2.2601	Generator Loss: 0.1397
Epoch 2 of 2	Steps: 2400	Discriminator Loss: 0.4738	Generator Loss: 2.4062
Epoch 2 of 2	Steps: 2500	Discriminator Loss: 0.6523	Generator Loss: 1.1895

Sample Output:




Epoch 2 of 2	Steps: 2600	Discriminator Loss: 1.1630	Generator Loss: 0.4774
Epoch 2 of 2	Steps: 2700	Discriminator Loss: 3.1739	Generator Loss: 0.0865
Epoch 2 of 2	Steps: 2800	Discriminator Loss: 3.0196	Generator Loss: 0.1123
Epoch 2 of 2	Steps: 2900	Discriminator Loss: 2.0650	Generator Loss: 0.1844
Epoch 2 of 2	Steps: 3000	Discriminator Loss: 1.2745	Generator Loss: 0.4915

Sample Output:




Epoch 2 of 2	Steps: 3100	Discriminator Loss: 2.5968	Generator Loss: 0.1318
Epoch 2 of 2	Steps: 3200	Discriminator Loss: 0.6263	Generator Loss: 1.4493
Epoch 2 of 2	Steps: 3300	Discriminator Loss: 0.6757	Generator Loss: 0.8448
Epoch 2 of 2	Steps: 3400	Discriminator Loss: 1.6737	Generator Loss: 3.3454
Epoch 2 of 2	Steps: 3500	Discriminator Loss: 0.1334	Generator Loss: 3.4634

Sample Output:




Epoch 2 of 2	Steps: 3600	Discriminator Loss: 2.7757	Generator Loss: 0.1358
Epoch 2 of 2	Steps: 3700	Discriminator Loss: 0.6593	Generator Loss: 1.7282
End of Epoch 2	Steps: 3750	Discriminator Loss: 1.0671	Generator Loss: 2.4526
Sample Output:			



1.3.4 Plot the losses for MNIST

```
In [85]: fig, ax = pyplot.subplots()
         losses = np.array(losses)
         pyplot.plot(losses.T[0], label='Discriminator', alpha=0.5)
         pyplot.plot(losses.T[1], label='Generator', alpha=0.5)
         pyplot.title("MNIST Training Losses")
         pyplot.legend()
```

```
Out [85]: <matplotlib.legend.Legend at 0x7fd9da111d68>
```



1.3.5 CelebA

Run your GANs on CelebA. It will take around 20 minutes on the average GPU to run one epoch. You can run the whole epoch or stop when it starts to generate realistic faces.

```
In [87]: batch_size = 32
         z_dim = 100
         learning_rate = 0.0002
         beta1 = 0.5
```

```

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

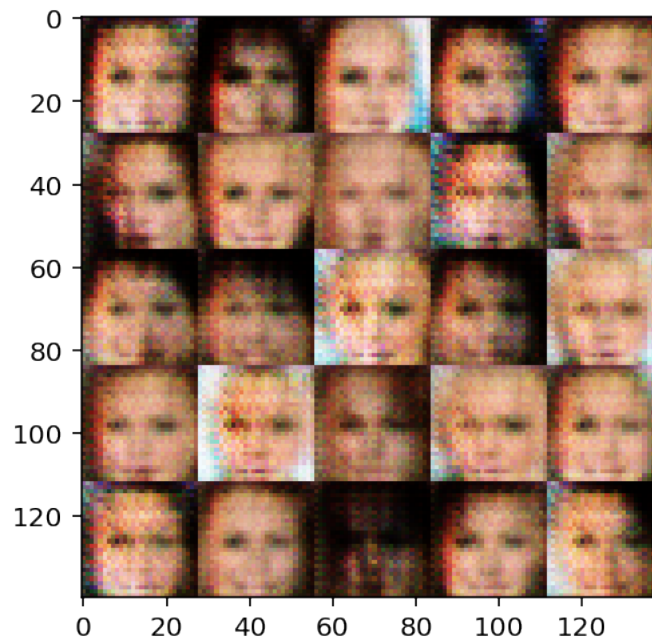
epochs = 1

celeba_dataset = helper.Dataset('celeba', glob(os.path.join(data_dir, 'img_align_celeba')))
with tf.Graph().as_default():
    train(epochs, batch_size, z_dim, learning_rate, beta1, celeba_dataset.get_batches(
        celeba_dataset.shape, celeba_dataset.image_mode)

```


Epoch 1 of 1	Steps: 100	Discriminator Loss: 0.1885	Generator Loss: 3.9552
Epoch 1 of 1	Steps: 200	Discriminator Loss: 0.5272	Generator Loss: 1.1354
Epoch 1 of 1	Steps: 300	Discriminator Loss: 1.1054	Generator Loss: 0.4556
Epoch 1 of 1	Steps: 400	Discriminator Loss: 0.7749	Generator Loss: 1.8894
Epoch 1 of 1	Steps: 500	Discriminator Loss: 0.9126	Generator Loss: 0.6933

Sample Output:



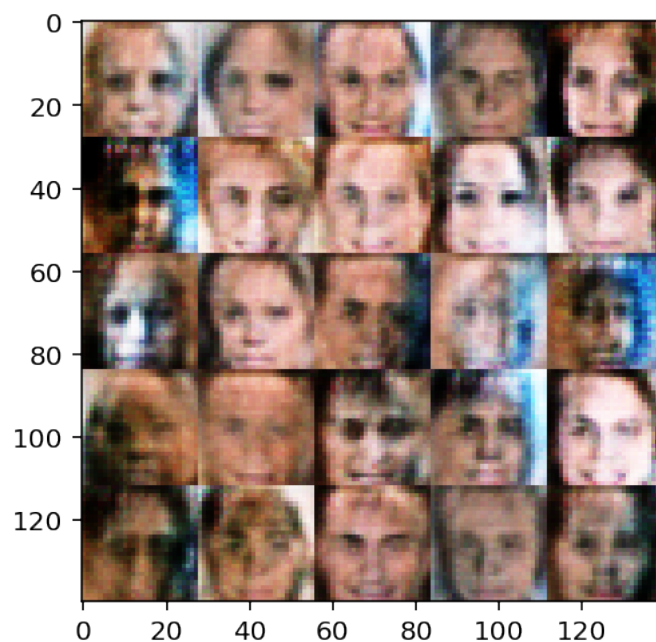
Epoch 1 of 1	Steps: 600	Discriminator Loss: 0.9069	Generator Loss: 0.8838
Epoch 1 of 1	Steps: 700	Discriminator Loss: 0.6950	Generator Loss: 1.2724
Epoch 1 of 1	Steps: 800	Discriminator Loss: 1.3792	Generator Loss: 0.3684
Epoch 1 of 1	Steps: 900	Discriminator Loss: 0.6706	Generator Loss: 1.7817
Epoch 1 of 1	Steps: 1000	Discriminator Loss: 1.2426	Generator Loss: 0.5058

Sample Output:




Epoch 1 of 1	Steps: 1100	Discriminator Loss: 0.6247	Generator Loss: 1.3149
Epoch 1 of 1	Steps: 1200	Discriminator Loss: 0.8420	Generator Loss: 1.4984
Epoch 1 of 1	Steps: 1300	Discriminator Loss: 0.5092	Generator Loss: 1.9321
Epoch 1 of 1	Steps: 1400	Discriminator Loss: 0.9239	Generator Loss: 0.8873
Epoch 1 of 1	Steps: 1500	Discriminator Loss: 0.3995	Generator Loss: 2.5064

Sample Output:




Epoch 1 of 1	Steps: 1600	Discriminator Loss: 0.6691	Generator Loss: 0.8472
Epoch 1 of 1	Steps: 1700	Discriminator Loss: 0.6267	Generator Loss: 1.0984
Epoch 1 of 1	Steps: 1800	Discriminator Loss: 0.9537	Generator Loss: 0.6033
Epoch 1 of 1	Steps: 1900	Discriminator Loss: 0.7679	Generator Loss: 1.1273
Epoch 1 of 1	Steps: 2000	Discriminator Loss: 1.1056	Generator Loss: 0.6643

Sample Output:




Epoch 1 of 1	Steps: 2100	Discriminator Loss: 0.7365	Generator Loss: 2.1508
Epoch 1 of 1	Steps: 2200	Discriminator Loss: 1.0225	Generator Loss: 2.7270
Epoch 1 of 1	Steps: 2300	Discriminator Loss: 0.7895	Generator Loss: 0.7321
Epoch 1 of 1	Steps: 2400	Discriminator Loss: 0.4528	Generator Loss: 1.3336
Epoch 1 of 1	Steps: 2500	Discriminator Loss: 0.3408	Generator Loss: 1.4557

Sample Output:



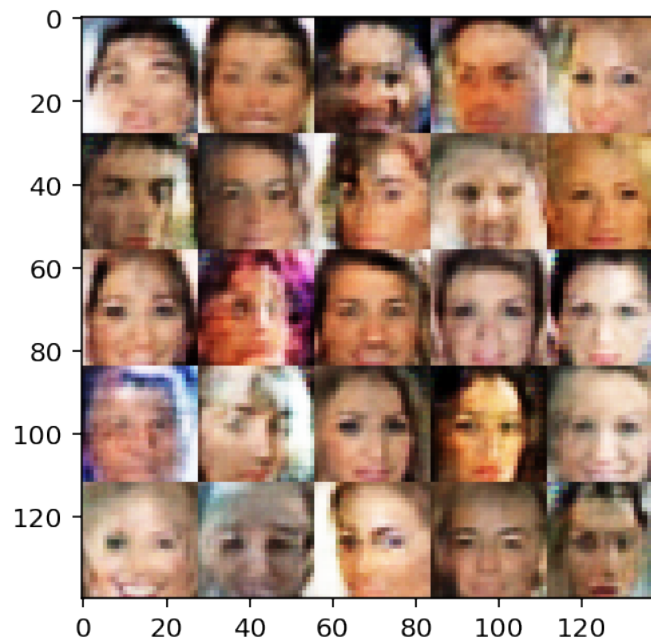
Epoch 1 of 1	Steps: 2600	Discriminator Loss: 0.7500	Generator Loss: 0.8244
Epoch 1 of 1	Steps: 2700	Discriminator Loss: 0.9716	Generator Loss: 0.6389
Epoch 1 of 1	Steps: 2800	Discriminator Loss: 4.3122	Generator Loss: 7.5529
Epoch 1 of 1	Steps: 2900	Discriminator Loss: 0.2864	Generator Loss: 3.0873
Epoch 1 of 1	Steps: 3000	Discriminator Loss: 0.3016	Generator Loss: 1.8224

Sample Output:



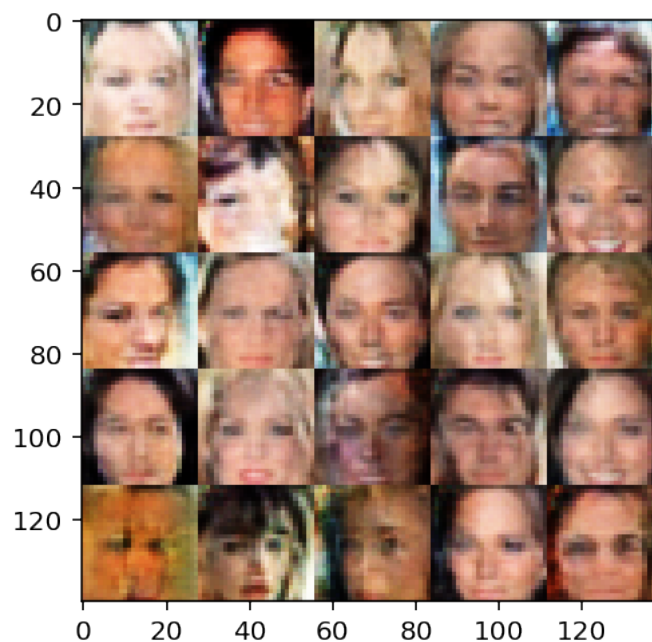
Epoch 1 of 1	Steps: 3100	Discriminator Loss: 1.1744	Generator Loss: 0.4826
Epoch 1 of 1	Steps: 3200	Discriminator Loss: 0.5239	Generator Loss: 1.6302
Epoch 1 of 1	Steps: 3300	Discriminator Loss: 0.4323	Generator Loss: 1.7159
Epoch 1 of 1	Steps: 3400	Discriminator Loss: 0.6450	Generator Loss: 1.2080
Epoch 1 of 1	Steps: 3500	Discriminator Loss: 0.6287	Generator Loss: 2.4208

Sample Output:




Epoch 1 of 1	Steps: 3600	Discriminator Loss: 0.4985	Generator Loss: 1.2989
Epoch 1 of 1	Steps: 3700	Discriminator Loss: 0.6421	Generator Loss: 1.1014
Epoch 1 of 1	Steps: 3800	Discriminator Loss: 1.1983	Generator Loss: 0.4689
Epoch 1 of 1	Steps: 3900	Discriminator Loss: 0.2586	Generator Loss: 3.9556
Epoch 1 of 1	Steps: 4000	Discriminator Loss: 0.4266	Generator Loss: 1.5190

Sample Output:




Epoch 1 of 1	Steps: 4100	Discriminator Loss: 0.8941	Generator Loss: 2.9652
Epoch 1 of 1	Steps: 4200	Discriminator Loss: 0.8230	Generator Loss: 0.8529
Epoch 1 of 1	Steps: 4300	Discriminator Loss: 0.6564	Generator Loss: 1.2432
Epoch 1 of 1	Steps: 4400	Discriminator Loss: 0.5035	Generator Loss: 1.1539
Epoch 1 of 1	Steps: 4500	Discriminator Loss: 1.0658	Generator Loss: 0.5603

Sample Output:



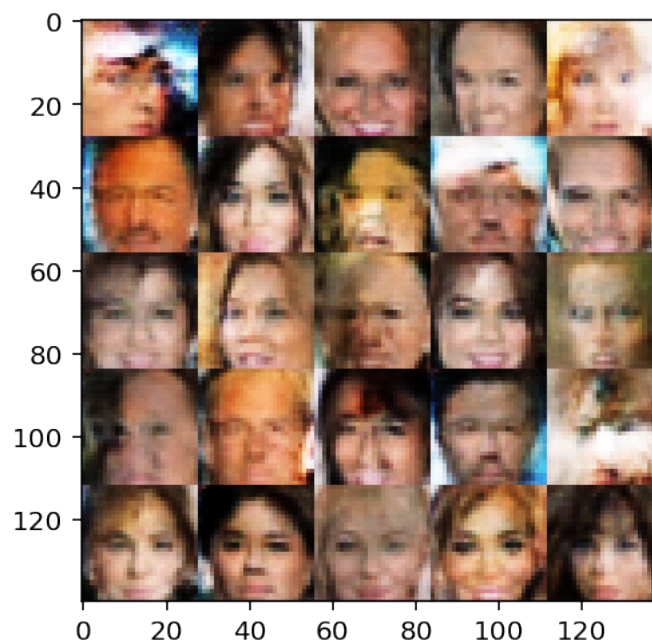
Epoch 1 of 1	Steps: 4600	Discriminator Loss: 0.5355	Generator Loss: 1.3598
Epoch 1 of 1	Steps: 4700	Discriminator Loss: 1.7372	Generator Loss: 4.7202
Epoch 1 of 1	Steps: 4800	Discriminator Loss: 1.1692	Generator Loss: 0.5474
Epoch 1 of 1	Steps: 4900	Discriminator Loss: 0.2905	Generator Loss: 2.5661
Epoch 1 of 1	Steps: 5000	Discriminator Loss: 0.8834	Generator Loss: 0.6470

Sample Output:




Epoch 1 of 1	Steps: 5100	Discriminator Loss: 1.0238	Generator Loss: 0.5904
Epoch 1 of 1	Steps: 5200	Discriminator Loss: 0.9232	Generator Loss: 0.6218
Epoch 1 of 1	Steps: 5300	Discriminator Loss: 0.2187	Generator Loss: 3.0423
Epoch 1 of 1	Steps: 5400	Discriminator Loss: 0.2330	Generator Loss: 2.2665
Epoch 1 of 1	Steps: 5500	Discriminator Loss: 2.0319	Generator Loss: 4.6782

Sample Output:




Epoch 1 of 1	Steps: 5600	Discriminator Loss: 0.2531	Generator Loss: 4.8671
Epoch 1 of 1	Steps: 5700	Discriminator Loss: 0.6378	Generator Loss: 1.0159
Epoch 1 of 1	Steps: 5800	Discriminator Loss: 1.0713	Generator Loss: 0.5014
Epoch 1 of 1	Steps: 5900	Discriminator Loss: 0.9492	Generator Loss: 0.6738
Epoch 1 of 1	Steps: 6000	Discriminator Loss: 1.0401	Generator Loss: 0.5536

Sample Output:



Epoch 1 of 1	Steps: 6100	Discriminator Loss: 1.4907	Generator Loss: 0.3084
Epoch 1 of 1	Steps: 6200	Discriminator Loss: 0.4664	Generator Loss: 2.0402
Epoch 1 of 1	Steps: 6300	Discriminator Loss: 0.5555	Generator Loss: 1.9956
End of Epoch 1	Steps: 6331	Discriminator Loss: 1.4238	Generator Loss: 0.3739


Sample Output:



1.3.6 Plot the losses for CelebA

```
In [88]: fig, ax = pyplot.subplots()
         losses = np.array(losses)
         pyplot.plot(losses.T[0], label='Discriminator', alpha=0.5)
         pyplot.plot(losses.T[1], label='Generator', alpha=0.5)
         pyplot.title("CelebA Training Losses")
         pyplot.legend()
```

```
Out [88]: <matplotlib.legend.Legend at 0x7fd9b3b0de80>
```



1.3.7 Submitting This Project

When submitting this project, make sure to run all the cells before saving the notebook. Save the notebook file as "dlnd_face_generation.ipynb" and save it as a HTML file under "File" -> "Download as". Include the "helper.py" and "problem_unittests.py" files in your submission.